

PHOTO: FLICKR IDLEMIND



More on Vulkan and SPIR-V: The future of high-performance graphics



Outline

- **Welcome**
 - Neil Trevett, Khronos President (NVIDIA)
- **Vulkan project overview**
 - Tom Olson, GL Next committee chair (ARM)
- **Vulkan applications**
 - Graham Sellers, Vulkan specification co-editor (AMD)
- **SPIR-V provisional specification**
 - John Kessenich, SPIR-V specification editor (LunarG)
- **Member progress reports and demos**
 - Various members
- **Next steps**
 - Tom again
- **Q&A / Panel discussion**

Khronos Connects Software to Silicon

Open Consortium creating
ROYALTY-FREE, OPEN STANDARD
APIs for hardware acceleration

Defining the roadmap for
low-level silicon interfaces
needed on every platform

Graphics, compute, rich media,
vision, sensor and camera
processing

Rigorous specifications AND
conformance tests for cross-
vendor portability

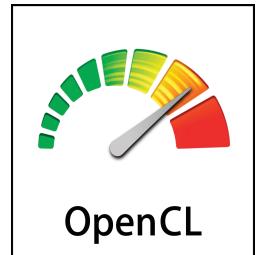
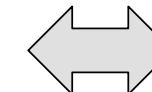
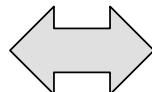
*Acceleration APIs
BY the Industry
FOR the Industry*



Well over a **BILLION** people use Khronos APIs
Every Day...

Khronos News at GDC

- Vulkan - next generation graphics API
 - Low overhead, high-efficiency graphics and compute on GPUs
 - Formerly discussed as Next Generation OpenGL Initiative
 - Technical overview and demos today - spec later this year
- SPIR-V - new shader IR supporting both graphics and compute constructs
 - Adopted by both Vulkan *and* OpenCL 2.1
 - Provisional specification available today





Vulkan Project Overview

Tom Olson, ARM
GDC, March 2015

Vulkan project history / status

- June to August 2014
 - Next Generation OpenGL project launch
 - Unprecedented commitment from all sectors of the industry
 - Project disclosure and call for participation at SIGGRAPH
- Since then...
 - Intense focus and a lot of hard work
 - Vulkan unveil at GDC 2015
- Status
 - Broad agreement on basic shape and semantics of the API
 - ‘alpha’ header file enabling experiments
 - API spec drafting is under way
 - SPIR-V spec drafting basically complete - provisional spec available

Vulkan vision and goals

- An open-standard, cross-platform 3D+compute API for the modern era
 - Compatibility break with OpenGL
 - Start from first principles
- Goals
 - Clean, modern architecture
 - Multi-thread / multicore-friendly
 - Greatly reduced CPU overhead
 - Architecture-neutral - full support for tile-based as well as direct renderers
 - Predictable performance through *explicit control*
 - Improved reliability and consistency between implementations

Vulkan in a nutshell

- **Modern architecture**
 - GL Context replaced by separate command buffers and dispatch queues
- **Thread-friendly**
 - Most object types are free-threaded
 - Application is responsible for synchronization
- **Low CPU overhead**
 - Error checking and dependency tracking are the application's job
 - Can opt in to a validation layer
- **Explicit control of when work is done**
 - Shader compilation and command generation happen at predictable times
 - Immutable state specified early to move driver work away from dispatch time

But first...



Continental

MEDIATEK

PIXAR

SONY

Apple

LUCASFILM
Ltd

EPIC
GAMES

Oculus VR™

BLIZZARD
ENTERTAINMENT

EA

VALVE

OXIDE

unity

HI CORP.

TRANSGAMING

RTT

intel

mobica

codeplay

SAMSUNG

NVIDIA.

ARM

BROADCOM.

VIVANTE

- Huge thanks to the whole Vulkan team!
 - New members are always welcome



THANKS
AMD!





Vulkan Applications

Graham Sellers, AMD
GDC, March 2015

Hi! I'm Graham Sellers

- AMD's OpenGL and Vulkan architect
- Represent AMD at OpenGL ARB
- Contributor of many OpenGL features and extensions
- Author of OpenGL SuperBible
- Spent the last year or so working on Vulkan
- I'm going to whip through a complete Vulkan application from startup to tear down

This is pseudo-code, not final API
We are still finalizing some details

<https://www.khronos.org/vulkan>

 @grahamsellers

Vulkan Application Startup

- Vulkan is represented by an “instance”
- Application can have multiple Vulkan instances
 - Each is independent
 - Eases middleware, subsystems, etc.
- Instance is owned by the loader
 - Aggregates drivers from multiple vendors
 - Responsible for discovery of GPUs
 - Makes multiple drivers look like one big driver supporting many GPUs

```
VK_APPLICATION_INFO appInfo = { ... };  
VK_ALLOC_CALLBACKS allocCb = { ... };  
VK_INSTANCE instance;  
  
vkCreateInstance(&appInfo, &allocCb, &instance);
```

Vulkan GPUs

- Vulkan instance creation takes:
 - Application info - tell Vulkan about your application
 - Allocation callbacks - Vulkan will allocate system memory using *your allocator*
- Once you have an instance, ask it about GPUs

```
uint32_t gpuCount;  
VK_PHYSICAL_GPU gpus[10];  
  
vkEnumerateGpus(instance, ARRAYSIZE(gpus), &gpuCount, gpus);
```

- Produces a list of GPUs, and a count
- GPUs can be from different vendors
 - Integrated + discrete
 - Multiple discrete GPUs in one system
 - Cross-GPU resource sharing and explicit multi-GPU support is in API

Vulkan GPU Info

- Query information about a GPU

```
VK_SOME_GPU_INFO_STRUCTURE info;  
uint32_t infoSize = sizeof(info);  
  
vkGetGpuInfo(gpu[0], VK_GPU_INFO_WHATEVER, &infoSize, &info);
```

- Lots of information available about GPU
 - Manufacturer, relative performance, memory sizes, queue types, etc.
- Cross-GPU compatibility query

```
VK_GPU_COMPATIBILITY_INFO compatInfo;  
  
vkGetMultiGpuCompatibility(gpuA, gpuB, &compatInfo);
```

- Compatibility info indicates
 - Full sharing, sharing of specific resources, or no compatibility at all

Vulkan Devices

- Construct a device instance from a GPU

```
VK_DEVICE_CREATE_INFO info = { ... };  
VK_DEVICE device;  
  
vkCreateDevice(gpu, &info, &device);
```

- Creation info contains information about:
 - Number and type of queues required
 - Which extensions you want to use
 - Extensions are ‘opt-in’ - cannot accidentally use an extension
 - Level of validation
 - Drivers generally will not include much, if any, error checking
 - Layers above can validate at various levels
 - Drivers may include multiple layers to validate vendor-specific behavior

Vulkan Queues

- Get queue handles from the device

```
VK_QUEUE queue;  
  
vkGetDeviceQueue(device, 0, 0, &queue);
```

- Queues are represented using two indices
 - Node ordinal
 - Node ordinal represents a “family” of queues, which are directly compatible
 - Queue index
 - Each queue family can have many queue instances
- Queues encapsulate
 - Functionality - graphics, compute, DMA
 - Scheduling - independently scheduled, asynchronous

Vulkan Command Buffers

- GPU commands are batched in command buffers

```
VK_CMD_BUFFER_CREATE_INFO info;  
VK_CMD_BUFFER cmdBuffer;  
  
vkCreateCommandBuffer(device, &info, &cmdBuffer);
```

- Create as many command buffers as you need
- Command buffer creation info includes
 - Which queue family you want to submit commands to (node ordinal)
 - Information about how aggressively drivers should optimize for GPU performance
 - etc.

Vulkan Commands

- Commands are inserted into command buffers

```
VK_CMD_BUFFER_BEGIN_INFO info = { ... };  
vkBeginCommandBuffer(cmdBuf, &info);  
  
vkCmdDoThisThing(cmdBuf, ...);  
vkCmdDoSomeOtherThing(cmdBuf, ...);  
  
vkEndCommandBuffer(cmdBuf);
```

- Driver heavy lifting happens here
 - Build many command buffers from many threads
 - Re-use command buffers
 - Spend time here optimizing work, not last minute right before draw
 - Big packages of immutable state make the workload less regardless

Vulkan Shaders

- Vulkan shaders are compiled up-front

```
VK_SHADER_CREATE_INFO info = { ... };  
VK_SHADER shader;  
  
vkCreateShader(device, &info, &shader);
```

- Shader creation info contains
 - Pointer to shader source
 - SPIR-V - portable, vendor-neutral, *open*, *extensible* shader binary
 - Other IRs could be supported through the same interfaces
 - Additional optional information
- Compile shaders from multiple threads
 - Driver will do as much work as it can right here

Vulkan Pipeline State

- Pipeline state is fully compiled

```
VK_GRAPHICS_PIPELINE_CREATE_INFO info = { ... };
VK_PIPELINE pipeline;

vkCreateGraphicsPipeline(device, &info, &pipeline);
```

- Creation info contains

- Compiled shaders
- Blend, depth, culling, stencil state, etc.
- List of states that need to be mutable

- Pipelines can be serialized and deserialized

```
uint32_t dataSize = DATA_SIZE;
void* data = malloc(DATA_SIZE);

vkStorePipeline(pipeline, &dataSize, data);
...
vkLoadPipeline(device, dataSize, data, &pipeline)
```

Vulkan Mutable State

- Some pipeline state is mutable or *dynamic*
- Represented by smaller state objects

```
VK_DYNAMIC_VP_STATE_CREATE_INFO vpInfo = { ... };  
VK_DYNAMIC_VP_STATE_OBJECT vpState;  
  
vkCreateDynamicViewportState(device, &vpInfo, &vpState);  
  
VK_DYNAMIC_DS_STATE_CREATE_INFO dsInfo = { ... };  
VK_DYNAMIC_DS_STATE dsState;  
  
vkCreateDynamicDepthStencilState(device, &dsInfo, &dsState);
```

Vulkan Resources

- Resources have a CPU and a GPU component
- CPU side is allocated using a `vkCreate*` function:

```
VK_IMAGE_CREATE_INFO imageInfo = { ... };
VK_IMAGE image;
vkCreateImage(device, &imageInfo, &image);

VK_BUFFER_CREATE_INFO bufferInfo = { ... };
VK_BUFFER buffer;
vkCreateBuffer(device, &bufferInfo, &buffer);
```

- It is the application's responsibility to allocate GPU memory for resources...

Vulkan GPU Memory

- **Query objects for their memory requirements**

```
VK_IMAGE_MEMORY_REQUIREMENTS reqs;  
size_t reqsSize = sizeof(reqs);  
  
vkGetObjectInfo(image,  
                 VK_INFO_TYPE_IMAGE_MEMORY_REQUIREMENTS,  
                 &reqsSize, &reqs);
```

- **Application allocates GPU memory**

```
VK_MEMORY_ALLOC_INFO memInfo = { ... };  
VK_GPU_MEMORY mem;  
vkAllocMemory(device, &memInfo, &mem);
```

- **Bind application-owned GPU memory to objects**

```
vkBindObjectMemory(image, 0, mem, 0);
```

Vulkan Descriptors

- Vulkan resources are represented by *descriptors*
- Descriptors are arranged in *sets*
- Sets are allocated from *pools*
- Each set has a layout, which is known at pipeline creation time
 - Layout is shared between sets and pipelines and must match
 - Layout represented by object, passed at pipeline create time
- Can switch pipelines which use sets of the same layout
- Many sets of various layouts are supported in one pipeline in a *chain*

```
vkCreateDescriptorPool(...);  
vkCreateDescriptorSetLayoutChain(...);  
vkCreateDescriptorSetLayout(...);  
vkAllocDescriptorSets(...);
```

Vulkan Render Passes

- Render passes represent logical phases of a frame
- Render passes are explicit objects

```
VK_RENDER_PASS_CREATE_INFO info = { ... };  
VK_RENDER_PASS renderPass;  
  
vkCreateRenderPass(device, &info, &renderPass);
```

- Render pass contains a lot of information about rendering
 - Layout and types of framebuffer attachments
 - What to do when the render pass begins and ends
 - The region of the framebuffer that the render pass may effect
- Vitally important information for tile-based and deferred renderers
 - ... but also very helpful for traditional forward-renderers!

Vulkan Drawing

- Draws are placed inside render passes
- Executed in the context of a command buffer

```
VK_RENDER_PASS_BEGIN beginInfo = { renderPass, ... };

vkCmdBeginRenderPass(cmdBuffer, &beginInfo);

vkCmdBindPipeline(cmdBuffer, VK_PIPELINE_BIND_POINT_GRAPHICS, pipeline);
vkCmdBindDescriptorSets(cmdBuffer, ...);
vkCmdDraw(cmdBuffer, 0, 100, 1, 0);

vkCmdEndRenderPass(cmdBuffer, renderPass);
```

- Pipelines, dynamic state objects and other resources bound to command buffers
- All draw types supported
 - Indexed and non-indexed, direct and (multi-)indirect, compute dispatches, etc.

Vulkan Synchronization

- Work is synchronized using event primitives

```
VK_EVENT_CREATE_INFO info = { ... };  
VK_EVENT event;  
  
vkCreateEvent(device, &info, &event);
```

- Events can be set, reset, queried and waited on

```
vkSetEvent(...);  
vkResetEvent(...);  
vkGetEventStatus(...);  
vkCmdSetEvent(...);  
vkCmdResetEvent(...);  
vkCmdWaitEvents(...);
```

- Command buffers can signal events as they complete execution

Vulkan Resource State

- Operations in command buffers are demarcated by pipeline barriers
- Barriers can wait on and signal events
- Barriers can transition resources from state to state
 - Renderable
 - Readable as texture
 - etc.

```
VK_IMAGE_MEMORY_BARRIER imageBarrier = { ... };
VK_PIPELINE_BARRIER barrier = { ..., 1, &imageBarrier };

vkCmdPipelineBarrier(cmdBuffer, &barrier);
```

- Drivers do not track state
 - Applications are responsible for state tracking
 - If you get it wrong, we will happily render garbage or crash
 - Validation layer will track state (slowly) and scream at you when you screw up

Vulkan Work Enqueue

- Work is executed on queues belonging to devices
- Completed command buffers are sent to queues for execution

```
VK_CMD_BUFFER commandBuffers[] = { cmdBuffer, ... };  
vkQueueSubmit(queue, 1, commandBuffers, fence);
```

- Queues own memory residency
 - Driver will not track memory residency for you

```
vkQueueAddMemReference(queue, mem);  
vkQueueRemoveMemReference(queue, mem);
```
- Queues can also signal and wait on semaphores for object ownership

```
vkQueueSignalSemaphore(queue, semaphore);  
vkQueueWaitSemaphore(queue, semaphore);
```

Vulkan Presentation

- Presentation is how we get images to the screen
- Displayable resource represented by a special kind of image
 - Bindable to framebuffers
 - Created by platform-specific modules called WSI (Window System Interface)
- Defining a small number (~2?) of WSI bindings
 - One for compositing systems where the compositor owns the displayable surface
 - One for systems that allow presentation of application-owned surfaces
- WSI also deals with things like:
 - Enumerating display devices and video modes
 - Going full screen
 - Controlling vsync
- Presentations enqueued along with command buffers

Vulkan Teardown

- Application responsible for object destruction
 - Must be correctly ordered
 - No reference counting
 - No implicit object lifetime
- Do not delete objects that are still in use!
- Most objects destroyed with:

```
vkDestroyObject(object);
```

- Some objects are “special”:

```
vkDestroyDevice(device);  
vkDestroyInstance(instance);
```

Vulkan AZDO?

- Vulkan is already PDCTZO (Pretty Darn Close to Zero Overhead)!
 - Very little validation unless you opt-in
 - You manage everything - virtually no driver funky business
 - Much better abstraction of the hardware - no complex mapping of API to silicon
- Submit the same command buffer many times
 - Amortized cost of building command buffer literally approaches zero
- Bindless
 - Debatable need - descriptor sets can be of arbitrary size
 - Explicit memory residency already in API
- Sparse
 - Yes
- MultiDrawIndirect
 - Yes
- Shader Draw Parameters
 - Yes

Vulkan Summary

- Vulkan is not “low level” - just a better abstraction of modern hardware
- Vulkan is very low overhead
 - Reduced CPU utilization means more cycles for your application
 - Explicit threading support means you can go wide without worrying about graphics APIs
 - Building command buffers once and submitting many times means low amortized cost
- Cross-platform, cross-vendor
 - Not tied to single OS (or OS version)
 - Not tied to single GPU family or vendor
 - Not tied to single architecture
 - Desktop + mobile, forward and deferred, tilers all first class citizens
- Open, extensible
 - Khronos is an open standards body
 - Collaboration from a wide cross-section of industry, IHVs + ISVs, games, CAD, AAA + casual
 - Full support for extensions, layering, debuggers, tools
 - SPIR-V fully documented - *write your own compiler!*



Thanks!

<https://www.khronos.org/vulkan>

 @grahamsellers



SPIR-V Provisional

GDC, San Jose
March 2015



Standard Portable Intermediate Representation

Goal:

- 1) Portable binary representation of shaders and compute kernels for GPUs and parallel computers
- 2) Target for OpenCL C/C++, GLSL, and other shader languages

*Enables compiler ecosystem for
more portable shaders*

Why use SPIR?

Without SPIR:

- Vendors shipping source
 - Risk IP leakage
- Limited Portability
 - No ISV control over front end
 - Different front end semantics per vendor
- Higher runtime compilation time

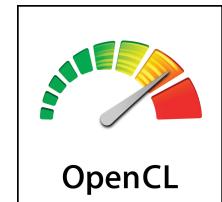
With SPIR:

- Ship a single binary
 - Requires tools to decipher; protecting IP
- Improved Portability
 - ISV can create their own front end tool chain
 - Multiple ISVs can share a common front end
- Reduced runtime compilation time
 - Some steps are offloaded

*Opportunity to unleash innovation:
Domain Specific Languages, C++ Compilers,*

What is SPIR-V?

- New intermediate language for input to Khronos graphics and compute APIs
 - Fully specified Khronos-defined standard
 - Can natively represent Khronos graphics and compute idioms
 - E.g., implicit derivatives with control-flow constraints
 - Memory and execution models for all GLSL and OpenCL high-level languages
- Core for Vulkan
 - The only language accepted by the API
 - Exposes machine model for Vulkan
 - Fully supports the GLSL/ESSL shader languages
 - Other shading languages easily target SPIR-V
- Core for OpenCL 2.1
 - Supports OpenCL 1.2, 2.0, 2.1 kernel languages



SPIR-V shader-language support

- Compiler chain split in two
 - Front end compiler emits SPIR-V portable binary IL, offline
 - SPIR-V IL is compiled to machine-specific binary by driver, online
- Front end NOT required in driver
 - Khronos working on offline language front ends

SPIR-V: A Deeper Look

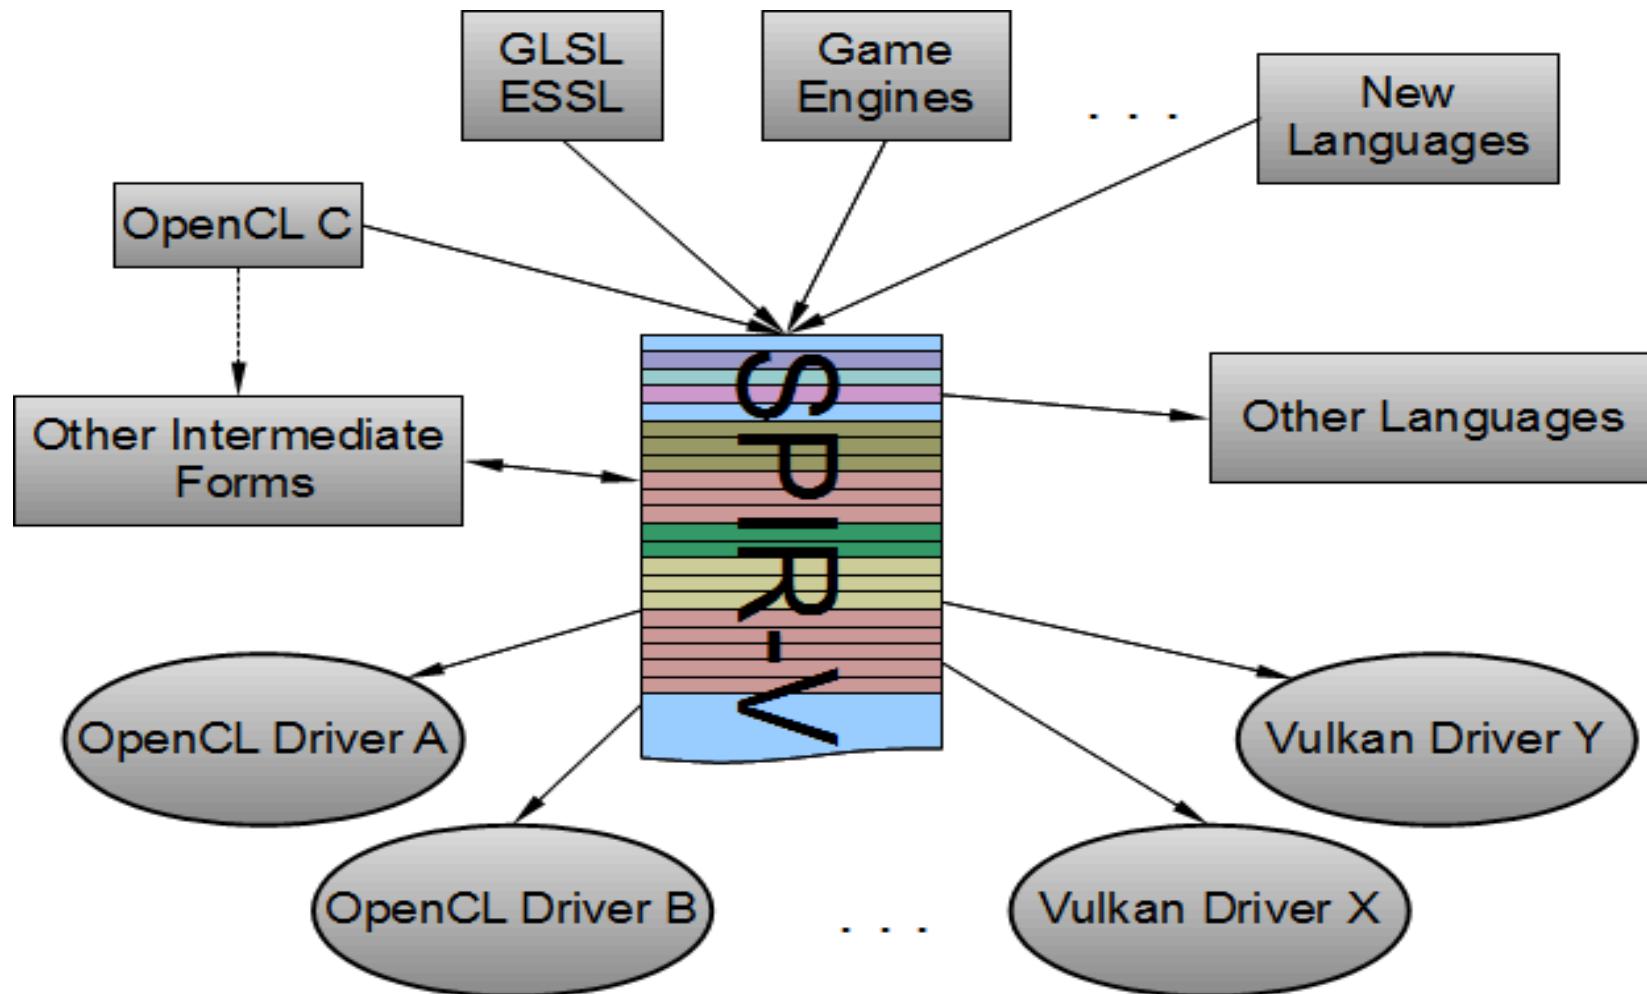
- A Binary Intermediate Language
 - A linear stream of words (32-bits)
- Functions inside a module contain a CFG (control-flow graph) of basic blocks
- Load/Store instructions are used to access declared variables
- Intermediate results are represented using single static-assignment (SSA)
- Data objects are represented logically, with hierarchical type information
 - e.g. No flattening of aggregates or assignment to physical registers
- Selectable addressing model
 - Allow usage of pointers, or dictate a memory model which is purely logical
- Can be easily extended
- Support debug information that can be safely stripped without changing the semantics of SPIR-V modules.

SPIR-V is a Binary Form

- Stream of words
- 32-bits wide
- Not a file format
 - This is the form passed through entry point
 - But, works well to start file with the magic number and directly store the stream
 - Deduce endianness from magic number

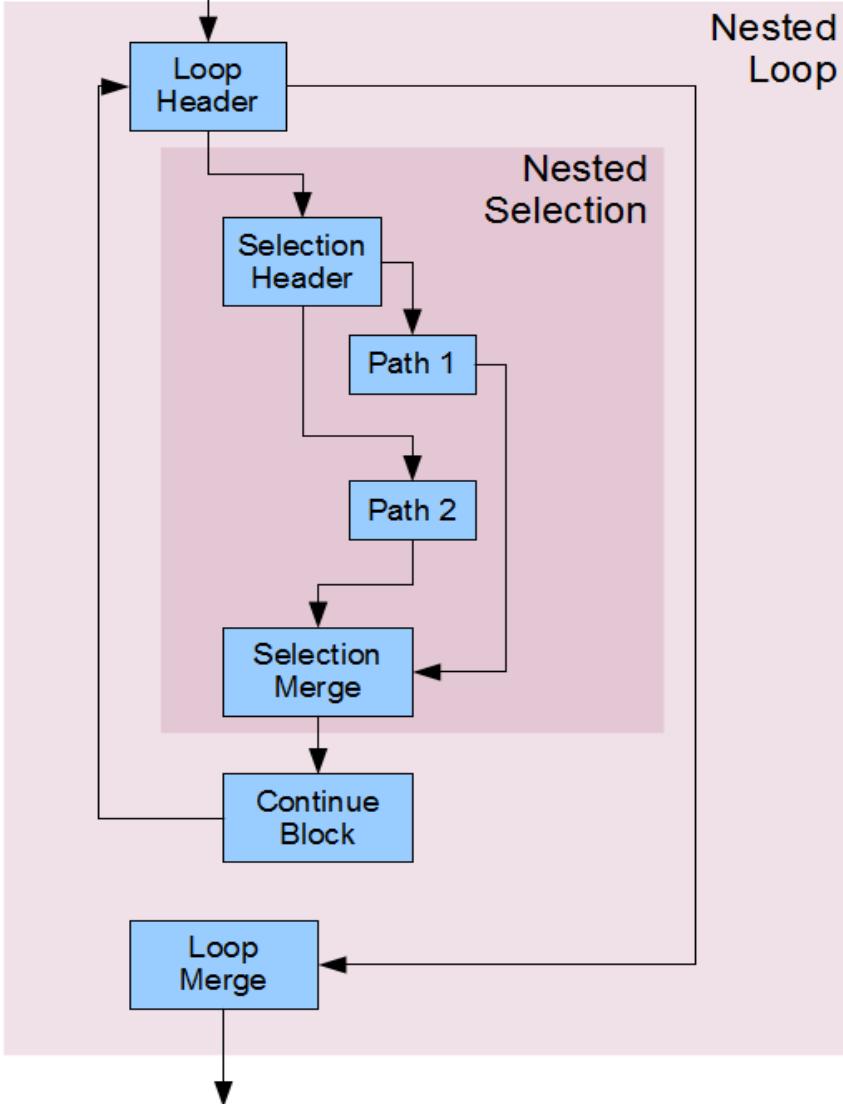
SPIR-V Magic #:	0x07230203
SPIR-V Version	99
Builder's Magic #:	0x051a00BB
<id> bound is	50
	0
OpMemoryModel	
	Logical
	GLSL450
OpEntryPoint	
	Fragment shader
	function <id> 4
OpTypeVoid	
	<id> is 2
OpTypeFunction	
	<id> is 3
	return type <id> is 2
OpFunction	
	Result Type <id> is 2
	Result <id> is 4
	0
	Function Type <id> is 3
	.
	.
	.

SPIR-V is a Common Intermediate Form



Structured Control Flow

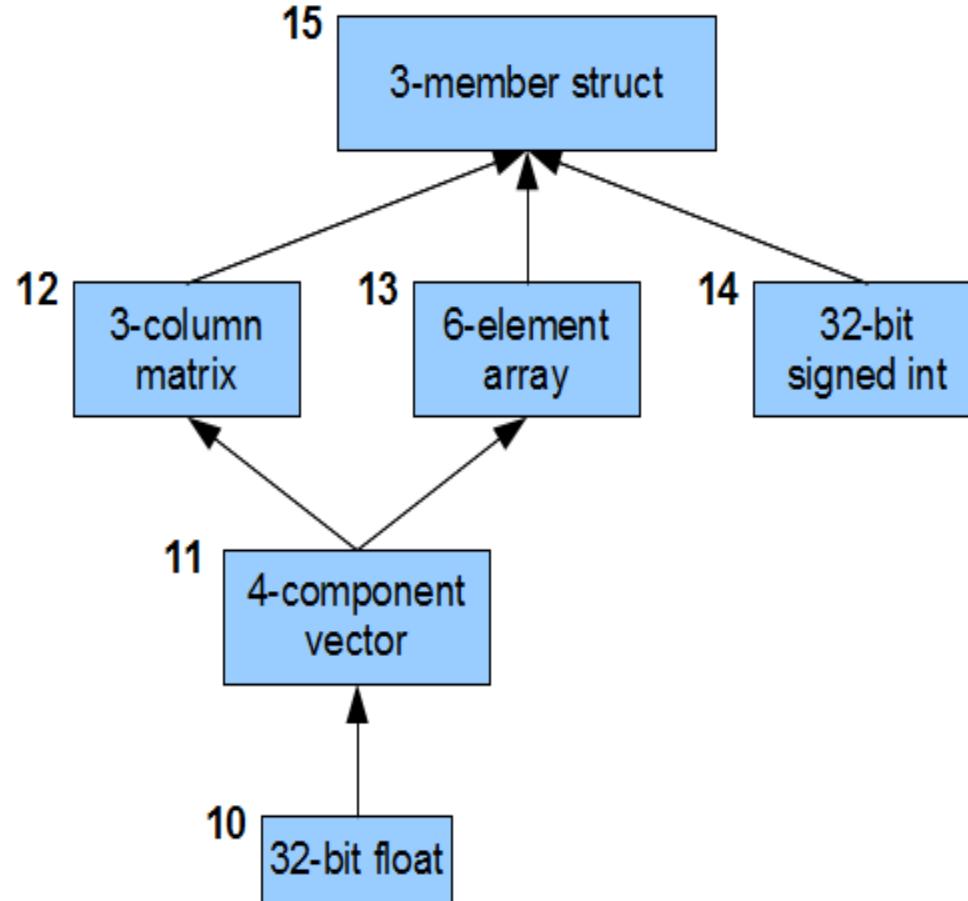
```
for (...) {  
  
    if (...)  
  
        ...  
    else  
        ...  
  
    ...  
  
}
```



```
11: Label  
...  
LoopMerge 12 NoControl  
BranchConditional 18 19 12  
19: Label  
22: ...  
SelectionMerge 24 NoControl  
BranchConditional 22 23 28  
23: Label  
...  
Branch 24  
28: Label  
...  
Branch 24  
  
24: Label  
...  
  
Branch 11  
  
12: Label
```

Hierarchical Types, Constants, and Objects

```
struct {  
    mat3x4;  
    vec4[6];  
    int;  
};
```



- 10: OpTypeFloat 32
- 11: OpTypeVector 10 4
- 12: OpTypeMatrix 11 3
- 13: OpTypeArray 11 6
- 14: OpTypeInt 32 1
- 15: OpTypeStruct 12 13 14

SPIR-V: A Deeper Look (Summary)

- A Binary Intermediate Language
 - A linear stream of words (32-bits)
- Functions inside a module contain a CFG (control-flow graph) of basic blocks
- Load/Store instructions are used to access declared variables
- Intermediate results are represented using single static-assignment (SSA)
- Data objects are represented logically, with hierarchical type information
 - e.g. No flattening of aggregates or assignment to physical registers
- Selectable addressing model
 - Allow usage of pointers, or dictate a memory model which is purely logical
- Can be easily extended
- Support debug information that can be safely stripped without changing the semantics of SPIR-V modules.

Call to Action



- Seeking feedback now on SPIR-V provisional
 - A Provisional specification, subject to change based on your feedback
 - Spec available at www.khronos.org/spir
 - Provide feedback at https://www.khronos.org/spir_feedback_forum
 - White paper <https://www.khronos.org/registry/spir-v/papers/WhitePaper.pdf>
- Innovate on the front end
 - New languages, abstractions
 - Target production quality Back ends
- Innovate on the back end
 - New target platforms: Multi core, Vector, VLIW...
 - Reuse production quality frontends
 - Other high-level languages and IRs/ILs
- Innovate on Tooling
 - Program analysis, optimization

PHOTO: FLICKR IDLEMIND



Member Progress Reports and Demos



 **Vulkan™** meets **Mali™**

Jesse Barker
Software Engineer, ARM

Vulkan investigations at ARM

- Prototype Vulkan driver for ARM® Mali™ Midgard GPU architecture
 - Intended to verify that Vulkan is a good fit to the architecture
 - Initial port on Arndale Octa (4+4 ARM Cortex™ A-15/7, Mali T-628 MP6)
- Caveats
 - Partial implementation – critical functions only, and some shortcuts
 - Built on top of an OpenGL ES / OpenCL HAL, not optimized for Vulkan

Experiment

- Draw call microbenchmark
 - 1000 meshes, 3 materials
 - Minimal state change between meshes
 - Measure CPU cycles in driver
 - Compare to OpenGL ES



Results

- For this test case, 79% reduction in CPU cycles spent in driver!

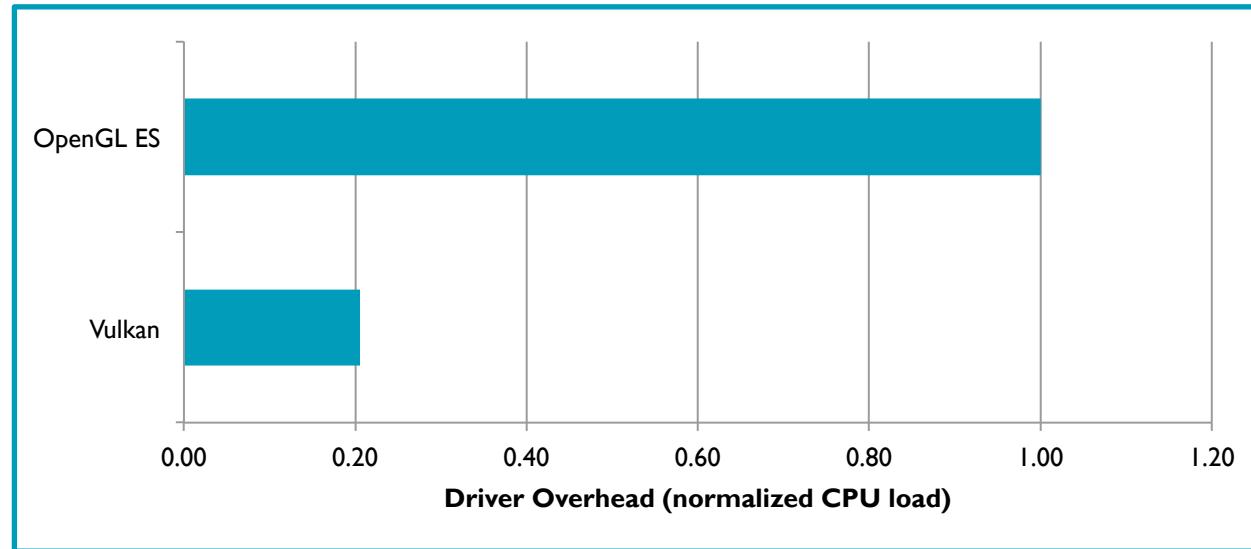


PHOTO: FLICKR IDLEMIND

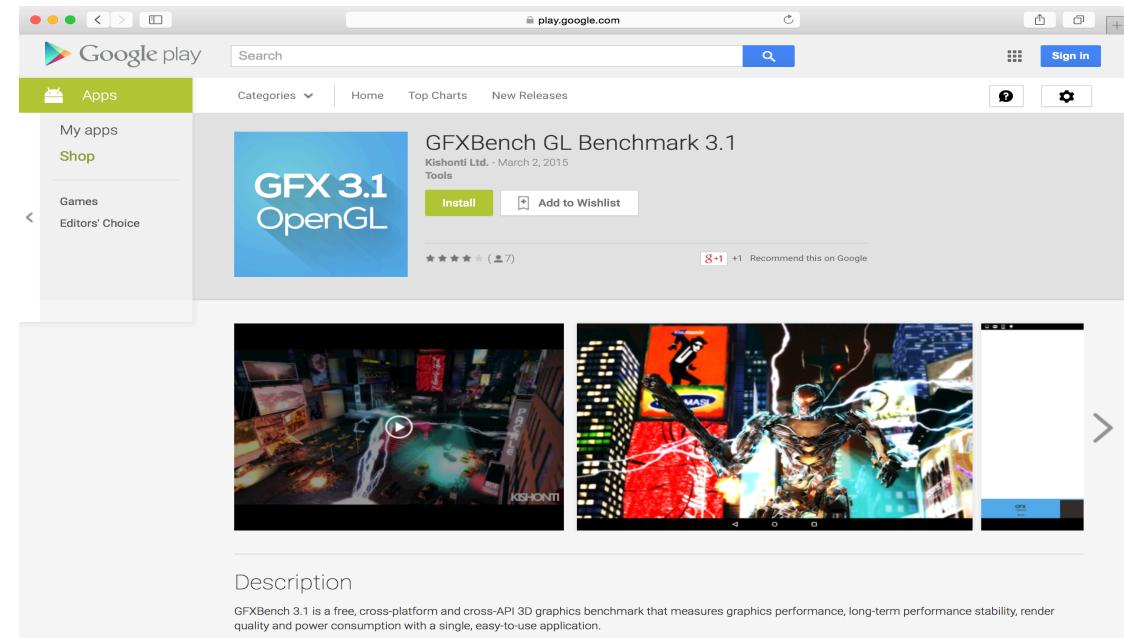


Kishonti Informatics



GFXBench 3.1

- Graphics benchmark for OpenGL ES 3.1
- In Google Play store this week!
- Adds compute shaders and new high-precision low level tests



GFXBench 4.0

- Graphics benchmark to showcase OpenGL ES 3.1 with Android Extension Pack (AEP)
 - Outdoor car chase scene with adaptive tessellation, HDR rendering, physically-based materials, compute post effects, dynamic reflections and shadows
 - Also sports geometry shaders and ASTC texture compression
 - Public release soon

GFXBench 5.0

- Entirely new engine aimed at benchmarking low-level graphics APIs (Vulkan, DX12, Metal)
- Concept is a night outdoor scene with aliens
- Still in pre-alpha, but shows the most important concepts
- Is showcased running Vulkan at Intel's GDC booth

PHOTO: FLICKR IDLEMIND



Imagination Technologies



PHOTO: FLICKR IDLEMIND



Intel

K H R O N O STM
G R O U P

PHOTO: FLICKR IDLEMIND



NVIDIA

K H R O N O STM
G R O U P

PHOTO: FLICKR IDLEMIND



Valve

K H R O N O STM
G R O U P

GLAVE debugger

LunarG.com/Vulkan

Call to Action

- Give us feedback on Vulkan and SPIR
 - Links provided on Khronos forums
 - https://www.khronos.org/spir_v_feedback_forum
 - https://www.khronos.org/vulkan/vulkan_feedback_forum
- Any company or organization is welcome to join Khronos for a voice and a vote in any of these standards
 - www.khronos.org
- Watch this space!
 - Initial specs and implementations coming later this year



PHOTO: FLICKR IDLEMIND



Q&A / Panel Discussion

